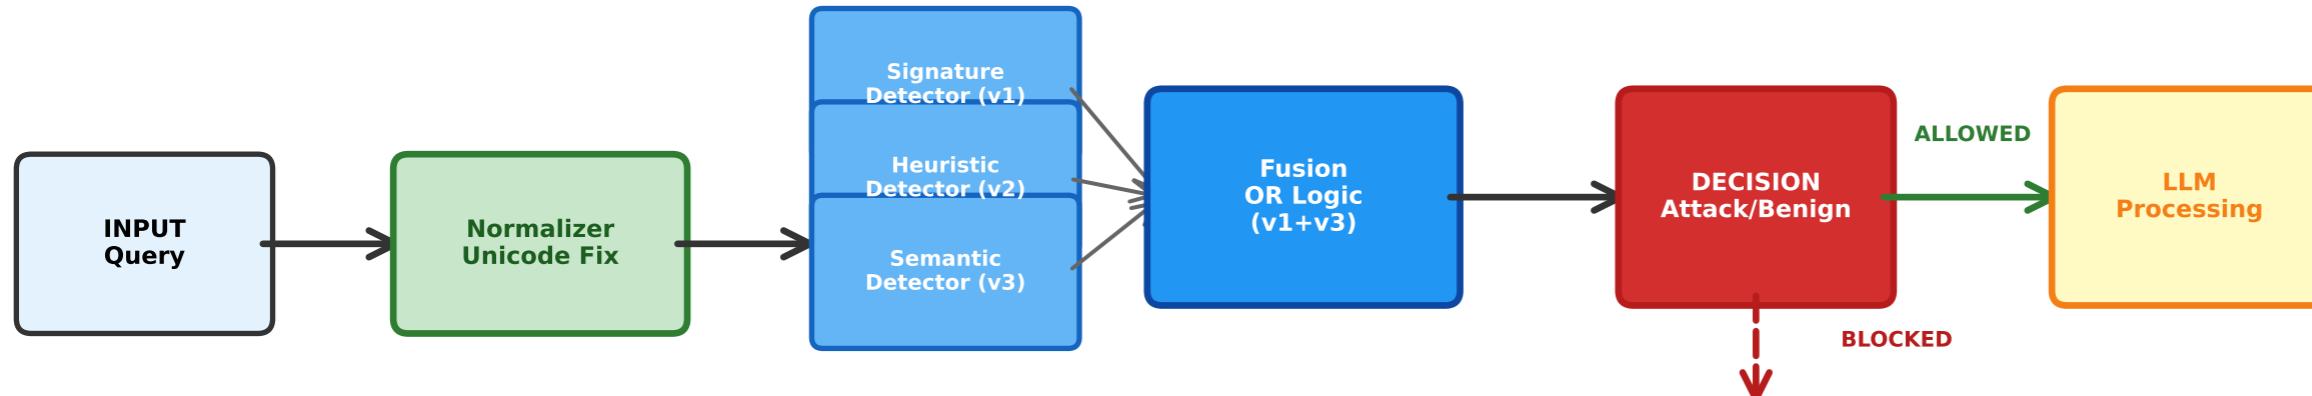
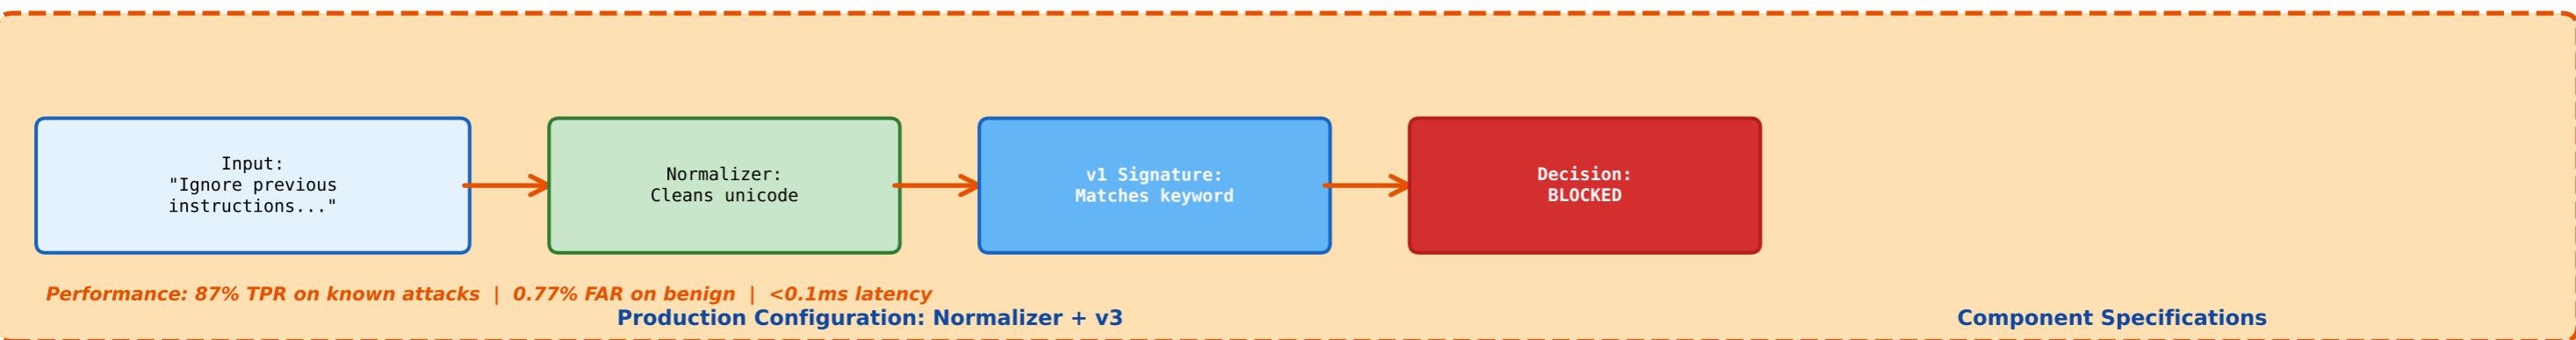


Prompt Injection Detection Pipeline Architecture

Input-Side Detection Before LLM Processing



Example: Attack Processing Flow



Component Specifications

True Positive Rate (TPR): 87% on known attacks
 False Alarm Rate (FAR): 0.77% on obfuscated benign
 Latency: <0.1ms per sample (CPU-only)
 Complexity: ~1,200 lines of code
 Deployment: Stateless, parallelizable
 Dependencies: None (pure Python)

Signature Detector (v1): Keyword matching

- 80% TPR, 0% FAR on Phase 1 attacks
- Catches: plain, delimiter, role confusion

Semantic Detector (v3): Pattern analysis

- 57% TPR, 0% FAR on Phase 1 attacks
- Catches: formatting, semantic anomalies

Fusion (OR Logic): v1 OR v3

- Combined: 87% TPR, 0% FAR

Key Design Principles

1. INPUT-SIDE DETECTION: Attacks blocked BEFORE reaching the LLM, preventing prompt injection at the source
2. NORMALIZER FIRST: Unicode/homoglyph normalization ensures consistent detection across obfuscation techniques
3. COMPLEMENTARY DETECTORS: v1 (signature) + v3 (semantic) catch different attack patterns through OR fusion
4. THRESHOLD-INVARIANT: Binary OR logic eliminates threshold tuning complexity in deployment
5. PRODUCTION-READY: <0.1ms latency, CPU-only, no external dependencies, stateless architecture

Legend:

□ Input

□ Normalizer

□ Detector

□ Fusion

□ Decision

□ LLM